

Nike Shoe Classifier - Project Documentation

Project Overview:

The Nike Shoe Classifier is an AI-based image classification model deployed as a FastAPI web service on Google Cloud Run. It allows users to upload images and classify them as:

- Nike
- Other Shoes
- Non-Shoe Items

The model is trained on a dataset containing Nike and non-Nike sneakers, along with non-shoe images to improve robustness.

Key Features:

- ❖ Multi-class classification (Nike, Other, Non-Shoe)
 - ❖ Deployed on Google Cloud Run for scalability
 - ❖ FastAPI with automatic Swagger UI (/docs)
 - ❖ Dockerized for easy deployment
 - ❖ Accessible via API for integration with web/mobile apps
-

Technologies Used:

Technology	Purpose
Python 3.9	Programming language
TensorFlow/Keras	Model training & inference
FastAPI	Web API framework
Uvicorn	ASGI server to run FastAPI
Docker	Containerization for deployment
Google Cloud Run	Serverless hosting
Swagger UI	Auto-generated API documentation
Google Container Registry	Stores Docker images

Project Structure:

```
/NikeClassifier
├── /app
│   ├── model.py      # Model loading & prediction functions
│   ├── __init__.py   # Package initialization
│   ├── main.py       # FastAPI application
│   ├── utils.py      # Helper functions (image preprocessing)
│   └── nike_final_model.keras #Model used for predictions
├── Dockerfile        # Docker setup for deployment
├── /Notebooks
│   └── Nikedifferentapproach.ipynb #Notebook that contains the code for the model
├── /Documentation
│   └── Documentation.docx #Documentation of the project
├── requirements.txt  # Python dependencies
└── run.sh           # Shell script for building & running Docker container
```

AI Model Documentation:

Project Overview

- **Problem Statement:** Develop an AI model that accurately classifies sneaker images into **Nike, Other Shoes, and Non-Shoe Items**.
- **Project Goals:** Improve sneaker classification accuracy for retail and authentication purposes.
- **Project Scope:** Limited to sneaker images; does not classify other apparel or brands outside the dataset.

Data Description

Data Sources

The dataset was compiled from multiple sources, including:

- **Kaggle:** Publicly available sneaker image datasets.
- **Web Scraping:** Scraped images from Nike and other shoe retailer websites.
- **Internal Datasets:** Manually collected images for better classification performance.
- **Open-source Repositories:** Leveraged sneaker classification datasets from research sources.

Data Preprocessing:

- Image resizing to **224x224** pixels
- Normalization (rescaling pixel values between 0-1)
- Data augmentation (flipping, zoom, rotation, brightness adjustment)

Data Split:

- Training set: **80%**
- Validation set: **20%**
- Testing set: **A New Dataset which is not trained by the model**

Data Quality Issues:

- Potential biases towards Nike branding features.
- Limited representation of rare sneaker models.

Model Details

Model Architecture:

- Base Model: MobileNetV2 (pre-trained on ImageNet)
- Batch Normalization: Normalizes activations to improve stability
- Global Average Pooling: Reduces dimensions while retaining information
- Fully Connected Layers:
 - 512 neurons, ReLU activation, L2 regularization (0.002)
 - 256 neurons, ReLU activation, L2 regularization (0.002)
- Dropout (0.4): Prevents overfitting
- Softmax Output Layer: Classifies into 3 categories (Nike, Other Shoes, Non-Shoe Items)
- Hyperparameters:
 - Optimizer: Adam (learning rate = 0.0003, decay applied) Loss Function: Categorical Crossentropy
 - Batch Size: 32
 - Epochs: 20 (initial), +10 fine-tuning

```
# Define the multi-class classification model
model = models.Sequential([
    base_model,
    layers.BatchNormalization(),
    layers.GlobalAveragePooling2D(),
    layers.Dense(512, activation='relu', kernel_regularizer=regularizers.l2(0.002)),
    layers.Dropout(0.4),
    layers.Dense(256, activation='relu', kernel_regularizer=regularizers.l2(0.002)),
    layers.Dropout(0.4),
    layers.Dense(3, activation='softmax') # Multi-class classification
])
```

Training Process

Training Overview

- The model is trained using transfer learning with MobileNetV2 as the base.
- The dataset is split into training, validation, and test sets.
- Data augmentation is applied to improve generalization and avoid overfitting.
- The training process is monitored using Early Stopping and ReduceLROnPlateau callbacks.
- Fine-tuning is performed after initial training to enhance feature learning.
- Validation Strategy:
 - Early stopping used to prevent overfitting.
 - ReduceLROnPlateau used to dynamically adjust learning rate.
- Training Metrics:
 - Accuracy, Loss, Learning Rate Adjustments.

```

: # Compile the model
optimizer = tf.keras.optimizers.Adam(learning_rate=0.0001)
model.compile(optimizer=optimizer,
               loss='categorical_crossentropy', # Multi-class loss function
               metrics=['accuracy'])

```

Callbacks Used

```

# Define callbacks
callbacks = [
    tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=8, restore_best_weights=True),
    tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.4, patience=4, min_lr=1e-6),
    tf.keras.callbacks.ModelCheckpoint("/Users/mani/nike_best_model.keras", monitor='val_loss', save_best_only=True),
    tf.keras.callbacks.TensorBoard(log_dir="./logs")
]

```

Initial Training

- The model is trained with frozen layers to extract generic features.
- It runs for **20 epochs**, with training resuming from epoch 17.

```

# Resume training from epoch 17
initial_epoch = 17
epochs = 20
history = model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=epochs,
    callbacks=callbacks,
    initial_epoch=initial_epoch
)

```

Fine-Tuning

- After initial training, all layers of the base model are unfrozen.
- The learning rate is reduced to **0.00005** for fine-tuning.
- Additional training is performed for **10 epochs** to refine feature learning.

```

# Fine-tuning after initial training
for layer in base_model.layers:
    layer.trainable = True # Unfreeze all layers
optimizer.learning_rate.assign(0.00005) # Reduce learning rate for fine-tuning

history = model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=10, # Additional fine-tuning epochs
    callbacks=callbacks
)

```

Bias Mitigation Strategies

To ensure fairness and reduce dataset bias, the following steps were taken:

- **Balanced Dataset Composition:** Efforts were made to collect equal representations of Nike, Other Shoes, and Non-Shoe images.
- **Augmentation Techniques:** Applied rotation, flipping, and color shifting to prevent over-reliance on specific visual patterns.
- **Oversampling & Undersampling:** Adjusted dataset ratios to balance underrepresented categories.
- **Bias Testing:** Evaluated predictions on diverse image sources to identify and correct skewed classifications.
- **Ongoing Monitoring:** Model retraining plans include continuous evaluation against new datasets to ensure bias reduction.

Key Insights from Classification Report

Nike Class

- **Precision:** 89% (Out of all Nike predictions, 89% were correct)
- **Recall:** 90% (Out of all actual Nike samples, 90% were identified correctly)

Other Class

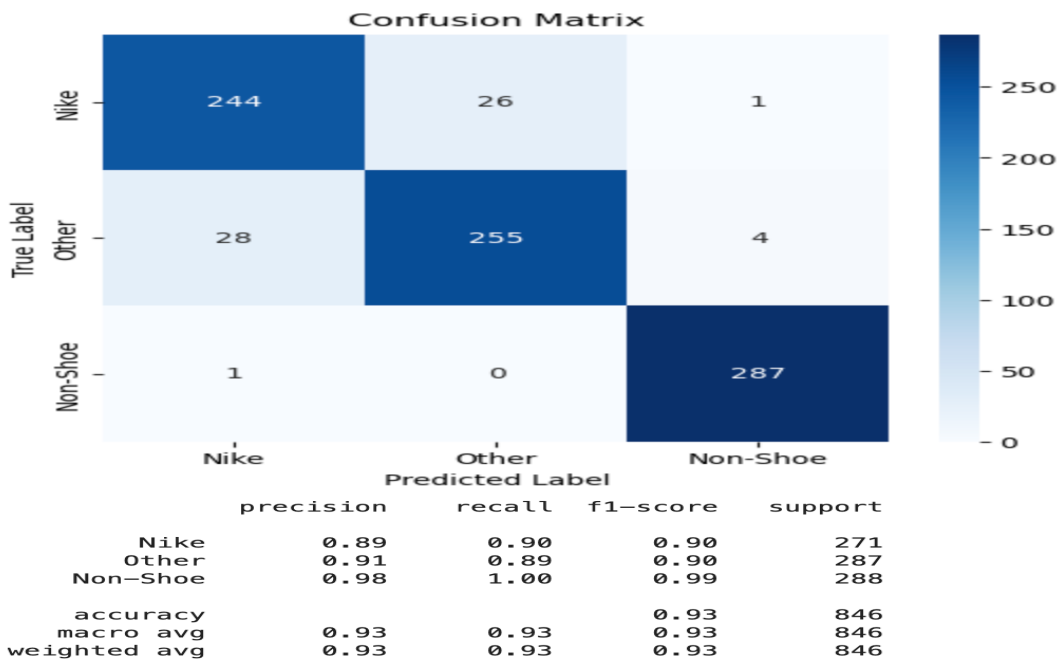

- **Precision:** 91%
- **Recall:** 89%

Non-Shoe Class

- Excellent classification with 98% Precision and 100% Recall, meaning nearly all Non-Shoe images were classified correctly.

Overall Model Performance

- **Accuracy:** 93% – A strong performance across all categories.
- **Balanced Precision, Recall, and F1-Scores** indicate reliable classification.
- **The model performs exceptionally well on distinguishing shoes from non-shoes** but has minor misclassification between Nike and Other brands.



Potential Biases and Limitations

Bias Analysis:

- More Nike sneaker images in training data could lead to overfitting.
- Dark-colored shoes occasionally misclassified due to limited contrast variations.

Limitations:

- Performance may degrade with unseen sneaker models not in training set.
- Requires clear images; low-resolution or occluded images reduce accuracy.

How the Model Works:

- User uploads an image via the API (/predict).
- Image is preprocessed (resized, normalized, etc.).
- Model predicts the class label.
- API returns a JSON response:

```
{
  "class": "Nike",
  "confidence": [[0.89, 0.07, 0.04]]
}
```

Cloud Deployed API:

- The API is hosted on Google Cloud Run with a public endpoint.
- To test with an image using a Python script:

```
import requests
url = https://nike-classifier-779080697580.us-central1.run.app/docs
files = {"file": open("test_shoe.jpg", "rb")}
response = requests.post(url, files=files)
print(response.json())
```

- Replace <https://nike-classifier-779080697580.us-central1.run.app/docs> with the actual deployed service URL.

How to Use the API:

1. Open Swagger UI for Interactive API Testing

Visit:

<https://nike-classifier-779080697580.us-central1.run.app/docs>

- Upload an image
- Click “Execute”
- View the classification result

2. Test API Using cURL:

Run the following command:

```
curl -X POST -F "file=@/path/to/image.jpg" \
https://nike-classifier-779080697580.us-central1.run.app
```

3. Use the API in a Python Script

```
import requests
url = https://nike-classifier-779080697580.us-central1.run.app
files = {"file": open("/path/to/image.jpg", "rb")}
response = requests.post(url, files=files)
print(response.json()) # {'class': 'Nike', 'confidence': [[0.89, 0.07, 0.04]]}
```

API Endpoints:

Endpoint	Method	Description
/	GET	Check API status
/docs	GET	Open Swagger UI
/predict	POST	Upload an image for classification

Example Request:

```
curl -X POST -F "file=@/path/to/image.jpg" \
  https://nike-classifier-xxxxxx.a.run.app/predict
```

Deployment Steps (Google Cloud Run):

- Set Google Cloud Project

```
gcloud config set project nikeidentifier
```

- Enable Required Services

```
gcloud services enable run.googleapis.com artifactregistry.googleapis.com
```

- Build and Push Docker Image

```
docker buildx build --platform linux/amd64 -t gcr.io/nikeidentifier/nike-classifier --push .
```

- Deploy to Cloud Run

```
gcloud run deploy nike-classifier \
  --image gcr.io/nikeidentifier/nike-classifier \
  --platform managed \
  --region us-central1 \
  --allow-unauthenticated \
  --port 8080 \
  --memory 2Gi
```

Security & Authentication:

Currently, the API is publicly accessible. Future enhancements include:

- API Key Authentication
- OAuth 2.0 Integration

- Google Cloud IAM Role-Based Access Control
-

Monitoring & Logging:

- Google Cloud Logging for request tracking
 - Google Cloud Monitoring for uptime checks
 - Error handling with structured logging
-

Future Enhancements:

- Improve model accuracy using more diverse training data
 - Add support for real-time image processing via Cloud Storage
 - Secure API with authentication & rate limiting
 - Develop a mobile-friendly UI for uploading images
-

Summary:

- Nike Shoe Classifier is a serverless FastAPI service deployed on Google Cloud Run.
- The model classifies images into Nike, Other, and Non-Shoe categories.
- The API is fully dockerized, scalable, and supports automated deployments.
- Swagger UI (/docs) allows easy testing of the API.

Project Status: Fully Deployed & Ready for Use!